

Special Issue on CAD/Graphics 2017

3D shape segmentation via shape fully convolutional networks

Pengyu Wang^{a,1}, Yuan Gan^{a,1}, Panpan Shui^a, Fenggen Yu^a, Yan Zhang^{a,*}, Songle Chen^b, Zhengxing Sun^a^a State Key Lab for Novel Software Technology, Nanjing University, Nanjing 210023, China^b Key Lab of Broadband Wireless Communication and Sensor Network Technology of Ministry of Education, Nanjing University of Posts and Telecommunications, Nanjing 210003, China

ARTICLE INFO

Article history:

Available online 9 August 2018

ABSTRACT

We design a novel fully convolutional network architecture for shapes, denoted by *shape fully convolutional networks (SFCN)*. 3D shapes are represented as graph structures in the SFCN architecture, based on novel *graph convolution and pooling operations*, which are similar to convolution and pooling operations used on images. Meanwhile, to build our SFCN architecture in the original image segmentation fully convolutional network (FCN) architecture, we also design and implement a *generating operation* with bridging function. This ensures that the convolution and pooling operation we have designed can be successfully applied in the original FCN architecture. In this paper, we also present a new shape segmentation approach based on SFCN. Furthermore, we allow more general and challenging input, such as *mixed datasets of different categories of shapes* which can prove the ability of our generalisation. In our approach, SFCNs are trained triangles-to-triangles by using three low-level geometric features as input. Finally, the feature voting-based multi-label graph cuts is adopted to optimise the segmentation results obtained by SFCN prediction. The experiment results show that our method can effectively learn and predict mixed shape datasets of either similar or different characteristics, and achieve excellent segmentation results.

© 2018 Elsevier Ltd. All rights reserved.

1. Introduction

Shape segmentation aims to divide a 3D shape into meaningful parts, and to reveal its internal structure. This is the basis and prerequisite to explore the inherent characteristics of the shape. The results obtained from shape segmentation can be applied to various fields of computer graphics, such as shape editing [1], deformation [2], and modelling [3]. Shape segmentation, therefore, has become a research hotspot, yet difficulties in the fields of digital geometric model processing and instance modelling persist.

Convolutional networks have shown excellent performance in various image processing problems, such as image classification [4–6], and semantic segmentation [7–9]. With the emerging encouraging study results, many researchers have devoted their efforts to various deformation studies on CNNs, one of which is the fully convolutional network (FCN) [10]. This method can train end-to-end, pixels-to-pixels on semantic segmentation, with no re-

quirement over the size of the input image. Thus, it has become one of the key research topics in CNN networks.

Although FCNs can generate good results in image segmentation, we cannot directly apply it to 3D shape segmentation. This is mainly because image is a type of a static 2D array, which has a very standard data structure and regular neighbourhood relations. Therefore, convolution and pooling can be easily operated when processing FCNs. However, the data structure of a 3D shape is irregular, so it cannot be directly represented as the data structure of an image. As triangle meshes have no regular neighbourhood relations like image pixels, direct convolution and pooling operations on a 3D shape is difficult to fulfil.

Inspired by the FCN architecture in image segmentation, we design and implement a new FCN architecture that operates directly on 3D shapes by converting a 3D shape into a graph structure. Based on the FCN process of convolution and pooling operation on the image and existing methods of Graph Convolutional Neural Networks [11–13], we design a shape convolution and pooling operation, which can be applied directly to the 3D shape. Combined with the original FCN architecture, we build a new shape fully convolutional network architecture and name it *shape fully convolutional network (SFCN)*. Secondly, following the SFCN architecture mentioned above and the basic flow of image segmentation of FCN

DOI of original article: [10.1016/j.cag.2017.07.030](https://doi.org/10.1016/j.cag.2017.07.030)

* Corresponding author.

E-mail address: zhangyannju@nju.edu.cn (Y. Zhang).¹ Both authors have equally contributed to this work.




Fig. 1. Example of our shape segmentation results on one mixed shape dataset. The shapes on the left are part of the training set, and some segmentation results are shown on the right.

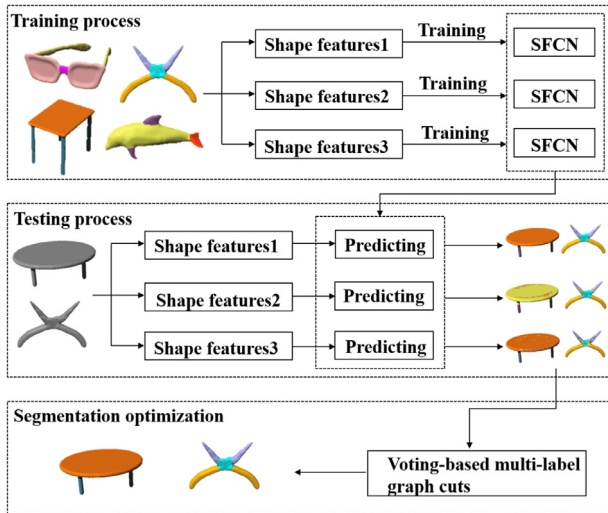


Fig. 2. The pipeline of our method. It may be divided into 3 stages: training process, using the SFCN architecture to train under three different features; testing process, predicting the test sets through the SFCN architecture; optimisation process, optimising the segmentation results by the voting-based multi-label graph cuts method to obtain the final segmentation results.

[10], we devise a novel trained triangles-to-triangles model for shape segmentation. Thirdly, for higher accuracy of segmentation, we use the multiple features of the shape to complete the training on the SFCN. Utilising the complementarity between features, and combined with the multi-label graph cuts method [14–16], we optimise the segmentation results obtained by the SFCN prediction. Our approach can realise the triangles-to-triangles learning and prediction with no requirements on the triangle numbers of the input shape. Furthermore, many experiments show that our segmentation results perform better than those of existing methods [17–19], especially when dealing with a large dataset. Finally, our method permits mixed dataset learning and prediction. Datasets of different categories are combined together in the test, and the accuracy of the segmentation results of different shapes decreases very little. As shown in Fig. 1, for a mixed shape dataset from COSEG [20] with several categories of shapes, part of the training set are displayed on the left, and some corresponding segmentation results are shown on the right. Fig. 2 shows the process of our method.

The main contributions of this paper include:

1. We design a fully convolutional network architecture for shapes, named shape fully convolutional network (SFCN), and is able to achieve effective convolution and pooling operations on the 3D shapes.

2. We present the shape segmentation and labelling based on SFCN. It can be triangles-to-triangles by three low-level geometric features, and outperforms the state-of-the-art shape segmentation.
3. Excellent segmentation results on training and predicting mixed datasets of different categories of shapes are achieved.

2. Related work

Fully convolutional networks. The fully convolutional network [10] proposed in 2015 is pioneering research which can effectively solve the problem of semantic image segmentation by pixel-level classification. Later, a great deal of research has emerged based on the FCN algorithms and achieved good results in various fields such as edge detection [21], depth regression [22], optical flow [23], simplifying sketches [24] and so on. However, the existing research on FCN is mainly restricted to image processing, largely because image has a standard data structure, easy for convolution, pooling and other operations.

Graphs CNN. Standard CNN cannot work directly on data which have graph structures. However, there are some previous works [12,13,25,26] that discuss how to train CNN on graphs. These works [13,25] use a spectral method by computing bases of the graph Laplacian matrix, and perform the convolution in the spectral domain. Niepert et al. [12] use a PATCHY-SAN architecture which maps from a graph representation onto a vector space representation by extracting locally connected regions from graphs.

Deep learning on 3D shapes. Recent works introduce various methods of deep architectures for 3D shapes. By using volumetric representation [27,28], standard CNN 2d convolution and pooling methods can be easily extended to 3D convolution and pooling. However, it leads to some problems which cannot be completely solved, such as data sparsity and the computation cost of 3D convolution. Kalogerakis et al. [29] proposed multi-view image-based Fully Convolutional Networks (FCNs) and surface-based CRFs to do the 3D mesh labelling by using RGB-D data as input. However, it may need a lot of time to compute the surface-pixel reference, and more viewpoints to maximally cover the shape surface. These methods [30,31] proposed different deep architectures which can be trained directly on point clouds to achieve points segmentation and other 3D points analysis. However, the point cloud structure sampled from the origin meshes may lack the intrinsic structure of the meshes. For better understanding the intrinsic structures, many works directly adapt neural networks to meshes. The first mesh CNN method GCNN [32] uses convolution filter in local geodesic polar coordinates. ACNN [33] used anisotropic heat kernels as filters. MoNet [34] considered a more generalised CNN architectures by using mixture model. Maron et al. [35] using a global seamless parameterisation for which the convolution operator is well defined but can only be used on sphere-like meshes. As




Fig. 3. Representation of different forms of data. (a) Image data representation; (b) shape data representation; (c) shape data represented as graph structure.

[36,37] show that intrinsic CNN methods can achieve better result on shape correspondence and descriptor learning with less training data. More details of these techniques can be found in [38].

Supervised methods for segmentation and labelling. When using the supervised method to train a collection of labelled 3D shapes, an advanced machine learning approach is used to complete the related training [17–20,39]. For example, Kalogerakis et al. [17] used Conditional Random Fields (CRF) to model and learn the sample example, in order to realise the component segmentation and labelling of 3D mesh shapes. Wang et al. [39] first projected 3D shapes to 2D space, and the labelling results in 2D space were then projected back to 3D shapes for mesh labelling. Xie et al. [19] use Extreme Learning Machines (ELM), which can be used to achieve consistent segmentation for unknown shapes. Guo et al. [18] and Kalogerakis et al. [29] applied deep architectures to produce the mesh segmentation.

Unsupervised segmentation and labelling. A lot of research methods can build a segmentation model [40–44] and achieve joint segmentation without any label information. There are predominantly two unsupervised learning methods: matching and clustering. Using the matching method, the matching relation between pair 3D shapes is obtained based on the similarity of a relative unit given by correlation calculation. The segmentation shape of matching shape is then established to realise the joint segmentation [40,45] of 3D shapes. By contrast, clustering methods analyse all the 3D shapes in the model set and cluster the consistent correlation units of 3D shapes into the same class. A segmentation model is then obtained and applied to consistent segmentation [42–44].

3. Shape fully convolution network

In the process of image semantic segmentation, it is mainly through operations such as convolution and pooling that fully convolution network architecture completes the image segmentation [10]. As analysed above, the regular data structure amongst the pixels of the image makes it easy to implement these operations. By analogy with images, triangles of the 3D shape can be seen as pixels on the image, but unlike pixels, triangles of the shape have no ordered sequence rules. Fig. 3(a) represents the regular pixels on the image, while Fig. 3(b) represents the irregular triangles on the shape. It is difficult to complete the convolution and pooling operation on the 3D shape like that of the image. Therefore, based on the characteristics of 3D shape data structure and analogous to the operation on the image, we need to design a new shape convolution and pooling operation.

As a 3D shape is mainly composed of triangles and the connections among them, we can use graph structure to describe it. Each 3D shape can be represented as $G = (V, E)$, with vertex $v \in V$ for a triangle and edge $e \in E \subset V \times V$ for the connection of adjacent triangles. The small triangle shown in 3(b) corresponds to the graph structure shown in Fig. 3(c). Based on the graph structure,

we design and implement a *shape convolution and pooling operation*, which will be detailed in the following section.

3.1. Convolution on shape

Convolution is one of the two key operations in the FCN architecture, allowing for locally receptive features to be highlighted in the input image. When convolution is applied to images, a receptive field (a square grid) moves over each image with a particular step size. The receptive field reads the feature values of the pixels for each channel once, and a patch of values is created for each channel. Since the pixels of an image have an implicit arrangement, a spatial order, the receptive fields always move from left to right and top to bottom. Analogous to the convolution operation on the image, therefore, we need to focus on the following two key ideas when employing convolution operation on the shapes:

1. *Determining the neighbourhood sets around each triangle for the convolution operation according to the size of the receptive field.*
2. *Determining the order of execution of the convolution operation on each neighbourhood set.*

For the first point, the convolution operation on the image is mainly based on the neighbourhood relationship between pixels. Accordingly, we need to construct locally connected neighbourhoods from the input shape. These neighbourhoods are generated efficiently and serve as the receptive fields of a convolutional architecture, permitting the architecture to learn shape representations effectively.

Shape has a neighbourhood relationship just like an image, but its irregularity restrains it to be directly represented and applied to the FCN learning. Yet when expressed as graph structures, the locally connected neighbourhoods of each triangle of the shape can be easily determined with various search strategies. In this paper, each triangle of the shape is viewed as a source node. We use the breadth-first search to expand its neighbourhood nodes on the constructed graph so as to obtain the neighbourhood sets of each triangle in the shape. Suppose the receptive field is set as K , the size of the neighbourhood set will be the same, including $K - 1$ neighbourhood nodes and a source node, all of which will be used for a follow-up convolution operation. Fig. 4(a) shows the graph structure of the shape, while Fig. 4(b) shows the neighbourhood sets of each source node (that is, each triangle on the 3D shape) we obtained with the method outlined above.

As for the second point, when performing the convolution operation on the image, it is easy to determine the order of the convolution operation according to the spatial arrangement of the pixels. However, it is rather difficult to determine the spatial orders among triangles on the 3D shape. A new strategy is thus needed to reasonably sort the elements in the neighbourhood sets. Sorting is to ensure that the elements in each neighbourhood set can be convolved by the same rules, so that the convolution operation can better activate features. For each node, all nodes in its neighbourhood set can be sorted by feature similarity (L2 similarity in the input feature space). Using this method, we can not only determine the order of the convolution operation of each neighbourhood set, but also ensure that the nodes in different sets have the same contribution regularity to their own source nodes in the convolution operation. The final convolution order for each neighbourhood set is shown in Fig. 4(c). As shown in Fig. 4(b), the execution order of the convolution operation of the neighbourhood set obtained from the source node is determined by calculating the input feature similarity.




Fig. 4. Convolution process on a shape. (a) Shape represented as a graph; (b) the neighbourhood nodes of different source nodes searched by a breadth-first search, among which 4, 6, 7 represent source nodes. The areas circled by orange, red and green dotted lines are the neighbourhood nodes searched by each source node. (c) Convolution order of each neighbourhood set. The ellipsis is for other neighbourhood sets not represented. (For interpretation of the references to colour in this figure legend, the reader is referred to the web version of this article.)

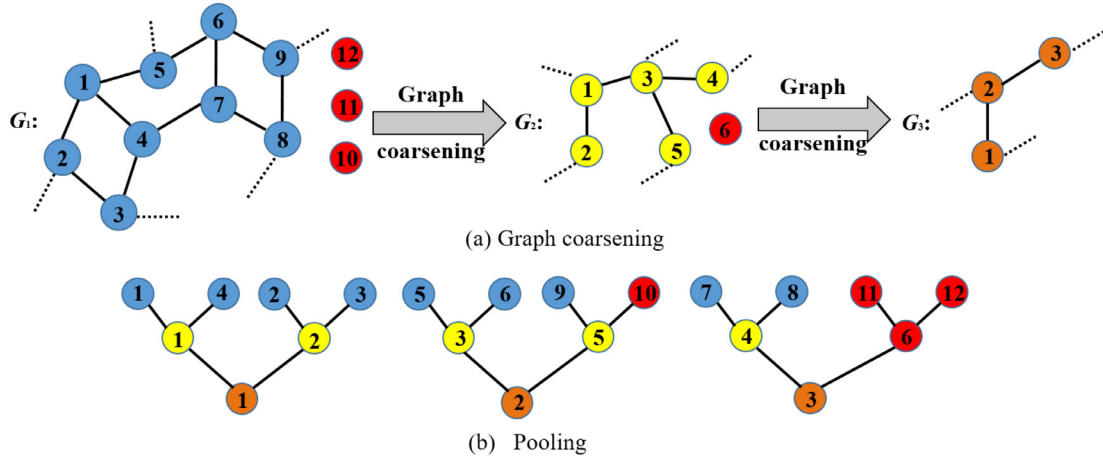


Fig. 5. Example of graph coarsening and pooling. (a) Graph coarsening process. Note: the original graph has 9 arbitrarily ordered vertices. For a pooling of size 4, two coarsenings of size 2 are needed. To ensure that in the coarsening process the balanced binary tree can be constructed, we need to add appropriate fake nodes through calculation, which are identified in red. After coarsening, the node order on G_3 is still arbitrary, yet it can be manually set. Then backstepping the coarsening process, we can determine the node order of G_2 and G_1 , and the corresponding relationship between nodes in each layer according to the order of G_3 . At that point the arrangement of vertices in G_1 permits a regular 1D pooling. (b) Pooling order. The nodes in the first layer (blue and red) represent the G_1 node order; the second layer (yellow and red) represents G_2 node order; the third layer (purple) represents the G_3 node order and the corresponding relationship between nodes in each layer. The red nodes are fake nodes, which are set to 0 in the pooling process, as we carry out maximum pooling. (For interpretation of the references to colour in this figure legend, the reader is referred to the web version of this article.)

3.2. Pooling on shape

Pooling is the other key operation in the FCN architecture. The pooling operation is utilised to compress the resolution of each feature map (the result of the convolution operation) in the spatial dimensions, leaving the number of feature maps unchanged. Applying a pooling operation across a feature map enables the algorithm to handle a growing number of feature maps and generalises the feature maps by resolution reduction. Common pooling operations are those of taking the average and maximum of receptive fields over the input map [11]. We share the same pooling operation on shape fully convolutional networks with the operation mentioned above. However, we need to address a key concern; that is, to determine the pooling operating order of the SFCN on the shape feature map.

Similar to the convolution operation, we cannot directly determine the pooling operation order on SFCN based on spatial relationships among the triangles of the 3D shape. Since the 3D shape has been expressed as a graph structure, we can determine the pooling operation order according to the operation of convolutional neural networks on the graph. In this paper, the pooling operation on the SFCN is computed by adopting the fast pooling of graphs [12,13].

The pooling method for graphs [12,13] coarsens the graph with the Graclus multi-level clustering algorithm [46]. Graph coarsening

aims to determine the new structure of the graph after pooling. We first present each shape as a graph structure, then we exploit the feature information on the triangles of the shape and the Graclus multi-level clustering algorithm [46] to complete shape coarsening; that is, to determine the new connection relationship of the shape feature map after pooling, which is shown in Fig. 5(a).

In the pooling process, traversing the feature map in a certain order according to the size of the receptive field is a key step to complete the calculation, namely, to determine the operation order of pooling. Following the method of pooling for the graph proposed by Defferrard et al. [13], the vertices of the input graph and its coarsened versions are irregularly arranged after graph coarsening. To define the pooling order, therefore, a balanced binary tree is built by adding fake code to sort the vertices. Lastly, the pooling operation of the graph is completed based on the nodes order and the use of a 1D signal pooling method. After shape coarsening, we apply the same approach in this study to determine the order of pooling operations on the shape fully convolutional networks architecture, as shown in Fig. 5(b).

4. Shape segmentation via shape fully convolution network

We design a novel shape segmentation method based on SFCN, analogous to the basic process of image semantic segmentation on FCN [10]. Firstly, we extract three categories of commonly used

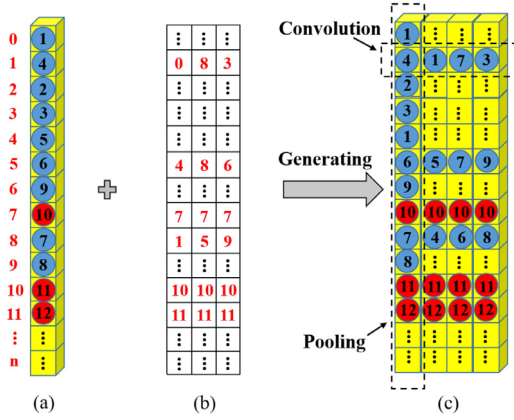


Fig. 6. Diagram of the generating layer. (a) The pooling order. A red number represents the offset, nodes in the blue circles represent nodes on the graph, nodes in the red circles are fake nodes. (b) Recorded neighbourhood set of each node. The numbers in the table represent the offset of each node after pooling sorting. (c) Data storage of the generating layer proposed in this paper, based on which our method can implement convolution operation by row and pooling operation by column. (For interpretation of the references to colour in this figure legend, the reader is referred to the web version of this article.)

geometric features of the shape as an input for SFCN training and learning. Secondly, based on the shape convolution and pooling operation proposed in Section 3 and the basic process of image semantic segmentation on FCN, we design a lattice structure suitable for 3D shape segmentation. Through training and learning of the network, we produce triangles-to-triangles labelling prediction. Finally, we introduce the optimisation process of shape segmentation.

4.1. Geometric feature extraction

Our approach is designed to complete the network training and learning based on some common and effective low-level features. In this paper, therefore, we extract three from the existing commonly used ones as the main features for the network training and learning. The three features include: average geodesic distance

(AGD) [47], shape context (SC) [48] and spin image (SI) [49]. These features describe the characteristics of each triangle on a shape from multiple perspectives well. We also found in the experiment that these three features are complementary, which will be analysed in detail in the experimental part.

4.2. Shape segmentation networks structure

As the convolution and pooling operations on a shape are different from those on an image, the FCN architecture originally used in image segmentation cannot be directly applied on the 3D shape. We modify the original FCN architecture according to the convolution and pooling characteristics, so that it can be conducted in shape segmentation.

Our training network is made up of four parts: *convolution*, *pooling*, *generating* and *deconvolution* layers, as shown in Fig. 7. The convolution layer corresponds to a feature extractor that transforms the input shape to multidimensional feature representation. The convolution operation is completed by the method proposed in Section 3.1. The pooling layer is used to reduce the feature vector of the convolution layer, and expand its receptive field to integrate feature points in the small neighbourhood into the new ones as output. The pooling operation is completed by the method proposed in Section 3.2.

As our convolution and pooling operations are designed for shapes, the original FCN architecture cannot be used directly. Compared with the original FCN's architecture, the architecture of the SFCN in this paper needs to record every neighbourhood set of each shape that participated in the convolution obtained in Section 3, as well as the convolution and pooling order of each shape. Thus, we add a *generating* layer in the original FCN architecture, whose diagram of concrete meaning is shown in Fig. 6.

Firstly, as shown in Fig. 6(a), we can calculate the pooling order between nodes of the graph by the shape pooling method proposed in Section 3.2 (these nodes are equivalent to the triangle of the shape). We store the nodes in the calculation order on the *generating* layer, as shown in Fig. 6(c). Fig. 6(a) gives the pooling order on a shape, where the figures represent the number of nodes (i.e. triangles).

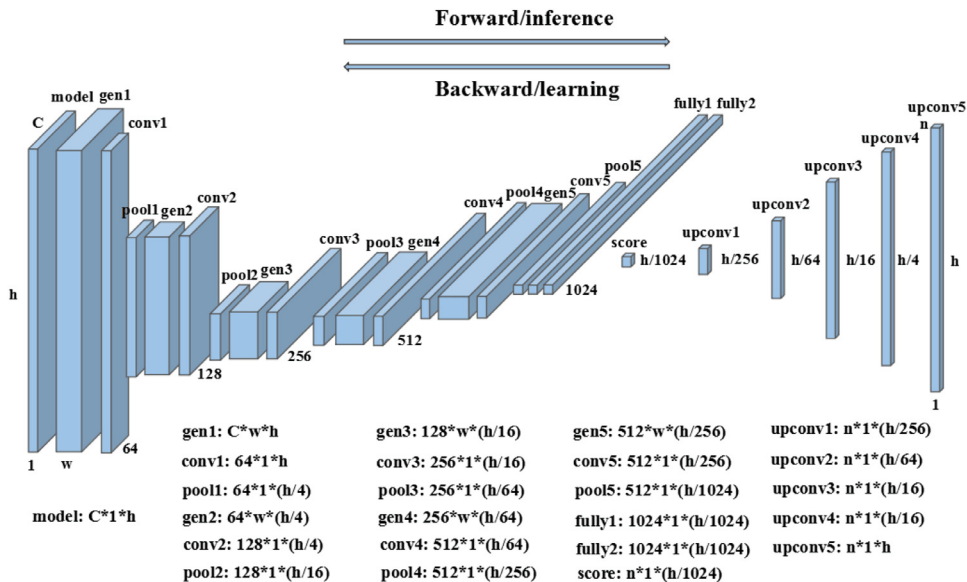


Fig. 7. The SFCN architecture we designed, mainly including 4 parts: convolution, pooling, generating, and deconvolution layers, which are represented by gen, con, pool and deconv respectively. Additionally, fully stands for fully connectivity and the numerical values reflect the dimension changes of each layer after calculation. C represents feature dimension, h represents the number of triangles of each shape, and n represents the shape segmentation labels.

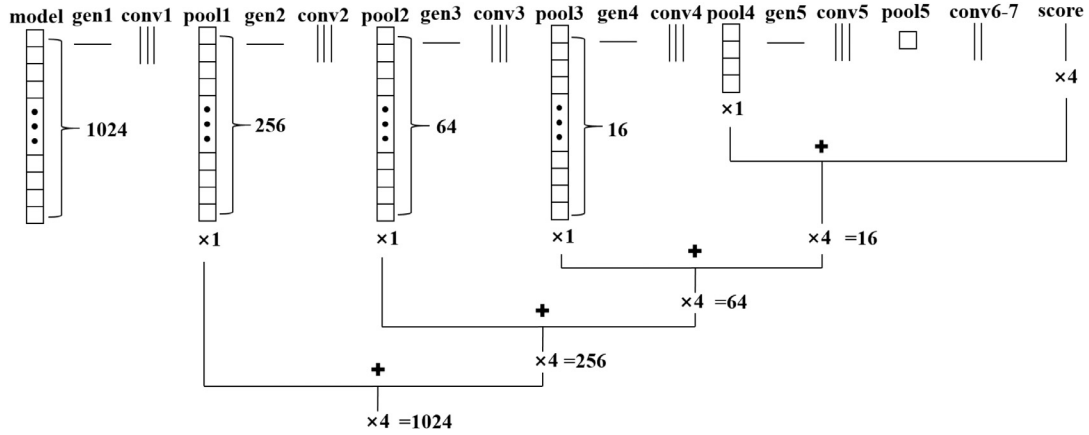


Fig. 8. The skip architecture for prediction. Our DAG nets learn to combine the coarse, high layer information with the fine, low layer one. We get the prediction results of the final layer by upsampling the score, then upsample again the prediction results combined with the final layer and the pool4 layer. After upsampling 5 times, we obtain the final prediction results, combining the final layer, pool5, pool4, pool3, pool2 and pool1 layer information, which then achieves triangles-to-triangles prediction.

Secondly, we need to record the neighbourhood sets of each node (i.e. each triangle of the shape) involved in the convolution computation, as shown in Fig. 6(b), where we store the neighbourhood sets by reading the offset in Fig. 6(a). After the nodes are sorted by column on the *generating layer*, we record their neighbourhood sets, in which the nodes are sorted according to the convolution order calculated in Section 3.1. As shown in Fig. 6(c), each row that is sequenced in the convolution order represents a neighbourhood set of a node, where the figures represent the number of nodes (i.e. triangles). By storing the data in this way, we can achieve the convolution operation by row and the pooling operation by column, as shown in Fig. 6(c). Another advantage of such storage is that, after pooling, the new neighbourhood set and the pooling order required by the next convolution layer of each new node can still be obtained and be applied to the next *generating layer* using the method in Section 3.2.

The deconvolution layer can be used to perform upsampling and density our graph representation. As mention above, we regard pooling as a graph coarsening by clustering 4 vertices, as shown in Fig. 7. Conversely, the upsampling operation can be regarded as a reverse process of the graph coarsening. We record how vertices are being matched before and after graph coarsening. Therefore it is easy to reuse the efficient deconvolution implementation based on the image-based FCN [10]. In this paper, the width of the convolution kernel in the deconvolution layer of the original FCN architecture is changed to 1 and the height is set to the size of the pooling we use, thereby obtaining the deconvolution layer of SFCN. The final output of the layer is a probability map, indicating the probability of each triangle on a shape that belongs to one of the predefined classes.

Based on the FCN architecture proposed by Long et al. [10], we design an SFCN architecture suitable for shape segmentation. Our convolution has five convolutional layers altogether, with each convolution layer having a generating layer before generating data for the next convolution and pooling, and followed by a pooling layer. Two fully connected layers are augmented at the end to impose class-specific projection. Corresponding to the pooling layer, there are five deconvolution layers, through which we can obtain the prediction probability of each triangle of the shape that belongs to each class. In the prediction process, we used the same skip architecture [10]. It can combine segmentation information from a deep, coarse layer with appearance information from a shallow, fine layer to produce accurate and detailed segmentations as the original FCN architecture. The specific process is shown in Fig. 8. The prediction probability of each layer can be obtained by adding

the results of a deconvolution layer and the corresponding results of the pooling layer after convolution, which also functions as the input for the next deconvolution layer. The number of rows will be repeated 5 times to return to the initial number of triangles, where the value of each channel is the probability of this class, realising the triangle level prediction. Another difference from the original FCN architecture is that, in order to normalise the data, we add a batch normalisation layer after the convolution operation of the first layer of the original FCN architecture, using the default implementation of the BN in the Caffe [50].

4.3. Shape segmentation optimisation

We train these three features separately using the network structure provided in Section 4.2. Given testing shapes, the SFCN produces the corresponding segmentation results under each feature, which can describe the segmentation of 3D shapes from different perspectives. Besides, due to the different input geometry features, there may be some differences among the predicted segmentation results of the same shape. To obtain the final segmentation results, we leverage the complementarity among features and the multi-label graph cuts method [14–16] to optimise the segmentation results of the testing shape. The final segmentation result is obtained mainly through the optimisation of the following formula.

$$E(I) = \sum_{u \in V} E_D(u, l_u) + \sum_{\{u, v\} \in E} E_S(u, v, l_u, l_v). \quad (1)$$

In this formula, l_u and l_v are labels of triangles u and v , data item $E_D(u, l_u)$ describes the energy consumption of triangle u marked as label l_u , and smoothing item E_S describes the energy consumption of neighbouring triangles marked using different labels.

The first item of the formula is optimised mainly based on the probability that triangle u is marked as label l_u . We predict the shape under the three features respectively, so the same triangle u will have its own prediction probability under each feature. In this paper, utilising the feature's complementarity, we vote the labelling results to get the final prediction probability, and serve its negative logarithm similar to the paper [18] as the first item of the multi-label graph cut. The second item in the formula smooths the segmentation results mainly through the calculation of the dihedral angle of the triangle and its neighbouring one. In this paper, the dihedral angle multiplied by the side length makes the second item of the formula to complete optimisation. Energy E is minimised by employing Multi-label Graph Cuts Optimisation [14–16],

Table 1
Labelling accuracy of large datasets.

	TOG (2011)	TOG (2013)	TOG (2015)	Ours
Chair	80.20%	91.20%	92.52%	93.11%
Vase	69.90%	85.60%	88.54%	88.91%

ToG (2011) is short for Sidi et al. [TOG 2011], ToG (2013) is short for Kim et al. [TOG 2013], ToG (2015) is short for Guo et al. (2015).

Table 2

Labelling accuracy of large datasets. Here are the results of the same test set for different numbers of training sets.

	Ours 25%	Ours 50%	Ours 75%	CGF (2014) 75%
Chair	93.43%	94.38%	95.91%	87.09%
Vase	88.04%	90.95%	91.17%	85.86%
Tele-alien	91.02%	92.76%	93.00%	83.23%

CGF (2014) is short for Xie et al. [CGF 2014].

Table 3

Labelling accuracy of mixed datasets of similar shapes.

	Airplane & Bird	Human & Teddy
Accuracy	90.04%	92.28%

through which we can obtain more accurate shape segmentation results.

5. Experimental results and discussion

Data. In this paper, deep learning is used to segment the shape. Therefore to verify the effectiveness of this method, we first tested 3 existing large datasets from COSEG [20], including the chair dataset of 400 shapes, the vase dataset of 300 shapes and the alien dataset of 200 shapes. The experiment of the mixed dataset was also carried out on the Princeton Segmentation Benchmark (PSB) [51] dataset. In addition, to confirm the robustness of the above method, we also selected 11 categories in the PSB [51] dataset for testing, each of them containing 20 shapes. For COSEG datasets, we chose the groundtruth as in the paper [20]. For PSB datasets, we chose the groundtruth as in the dissertation [17].

SFCN parameters. We trained by SGD with momentum, using momentum 0.9 and weight decay of $1e-3$. We chose ReLU as the activation function. Dropout used in the original classifier nets is included. The per-triangle, unnormalised softmax loss is a natural choice for segmenting shapes of any size, with which we trained our nets.

Computation time. We used two Intel(R) Xeon(R) CPU E5-2620 v3 @ 2.40 GHz with 12 cores and NVIDIA GeForce TITAN X GPU. In large datasets, when we used 50 shapes for training and the triangles of each shape range from 1000 to 5000, the training would take approximately 120 min (including the time to compute the input descriptors), and the test and optimisation of a shape took about 30 s.

Results. In the experiment, we used the classification accuracy proposed by Kalogerakis et al. [17] and Sidi et al. [41] for the quantitative evaluation of our method.

Firstly, to compare with the methods of [18] in the COSEG's large dataset, we randomly selected 50 shapes for training from the chair dataset and 20 shapes for training from the vase dataset. We compared our method with three shape segmentation methods [18,41,52]. The results are presented in Table 1. It should be noted that the results and data obtained by other methods come from [18]. The table shows that the results obtained by our method outperform the existing methods, thus proving that our method is effective.

Secondly, to verify the effectiveness of our method for large datasets, we randomly selected 25%, 50% and 75% of the shapes for training from each of the three large datasets from COSEG.

To verify whether the SFCN prediction accuracy becomes higher as the number of training sets increases, we used the same 25% shapes for testing in each experiment for each large dataset. In other words, the training sets gradually increased, while the test set remained the same. The results of this can be seen in Table 2. Because the work of [19] carried out a similar experiment under a 75% training condition with their method, we also make a comparison with their results in Table 2. It can be seen from the table that the results obtained by our method perform better than theirs. Furthermore, with the increase of the training set, the classification accuracy of our method grows steadily in the same test set. This shows that with the increase of the training set, both the learning ability and the generalisation ability of the network architecture become stronger, which also proves the effectiveness of the designed network in this paper.

Thirdly, we visualise the segmentation results obtained by using our method in the three large datasets of COSEG, as shown in Fig. 9. All the results shown are optimised and obtained using the 75% training set. The segmentation results appear visually accurate, which proves the effectiveness of this method.

Mixed dataset training and testing were performed as well. We respectively mixed airplane and bird, human and teddy, which are of similar classes. It must be noted that, unlike the method of [18] which merges similar class labels, we retained these different labels. So the body of the plane and bird have different labels; their wings as well. 12 shapes of each dataset were selected as the training set, and the remaining as the test set. Our approach was repeated three times to compute the average performance, the segmentation results of which are shown in Fig. 10. Fig. 10(a) is part of the result of the training set while Fig. 10(b) is part of the testing set. The classification accuracy of the two datasets is shown in Table 3, which suggests that we can obtain good segmentation results when mixing similar data together. Although the segmentation accuracy is not as high as training, the average is above 90%. Additionally, the basic segmentation is correct according to the visualisation of the results, proving that SPFCN network architecture is powerful in distinguishing features and learning.




Fig. 9. Results under large datasets.

Table 4
Labelling accuracy of mixed datasets of different shapes.

	Glasses & Plier	Chair (400) & Vase (300)	Glasses & Plier & Fish & Table
Accuracy	96.53%	87.71%	91.82%




Fig. 10. The segmentation results of mixed datasets of similar shapes. (a) Part of the shapes in the training set; (b) segmentation results of part of the shapes in the testing set.




Fig. 11. The segmentation results of mixed datasets of different shapes. (a) Part of the shapes in the training set; (b) segmentation results of part of the shapes in the testing set.

In addition, we mixed glasses and pliers, glasses, pliers, fish and tables of different categories. From the mixed datasets, we selected some data for training, with the remaining used for testing. Here 12 shapes of each dataset were selected as the training set, and the remaining as the test set. We also mixed two large datasets with 400 chairs and 300 vases. We selected 50% of the shapes for training from each dataset and the remaining 50% as the test set. Our approach was repeated three times to compute the average performance, the segmentation results of which are shown in Figs. 11 and 1. Fig. 11(a) is part of the result of the training set while Fig. 11(b) is part of the testing set. The classification accuracy of the three datasets is shown in Table 4, which also indicates that both the segmentation results and classification accuracy achieve impressive performance. In other words, this method can be used to train and predict any shape set, and proves once again that our SFCN architecture has good feature distinguishing ability and learning ability.

Lastly, as in the papers of [18], we separately trained several small datasets of PSB when $N = 6$ and $N = 12$ (i.e., SB6, SB12), in which N is the number of randomly selected shapes in each training process. For each category of dataset, as in the papers of [18], we repeated our approach five times to compute the average performance. The comparison results with the related methods are presented in Table 5. It should be noted that the results and data obtained by other methods come from [18]. As shown in Table 5, the results of several datasets of PSB obtained by our method perform much better in most categories than the existing methods [17–19,39,53,54], which proves the effectiveness of the method. On a few individual datasets, such as airplane, chair and table etc., our

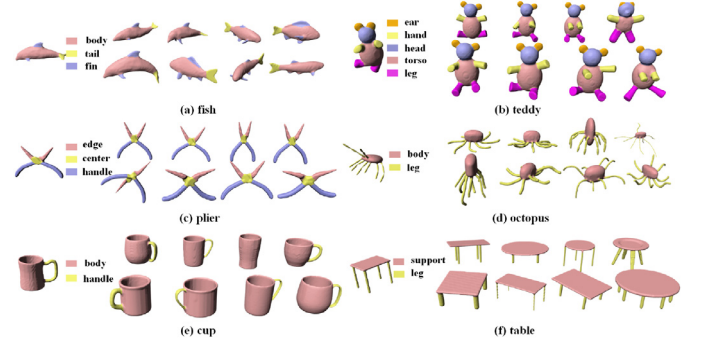


Fig. 12. More results of our method.

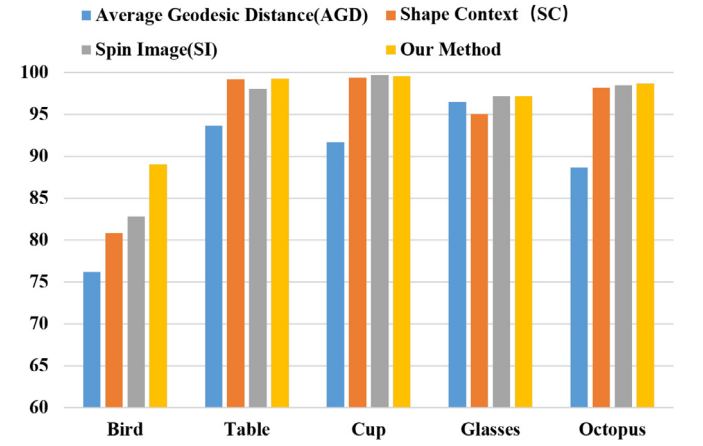


Fig. 13. The comparison of segmentation accuracy under different features.

results do not go beyond those of other methods, yet are very close to the best ones, which also serves to prove that our method is effective. Moreover, the segmentation effect gradually strengthens as the training data increases, which shows that the learning ability of the SFCN architecture is enhanced with the increase of training samples. We also visualise the segmentation results of several datasets of PSB including teddy, pliers and fish etc. on the condition of SB12, the optimised ones of which are shown in Fig. 12. Just like the large dataset, the results of the small one are visually accurate, which indicates that our method is feasible.

Feature sensibility. In this paper, we combine three features, average geodesic distance (AGD) [47], shape context (SC) [48] and spin image (SI) [49] to produce the shape segmentation. To verify the validity of this combination, we carried out a comparative experiment. The classification accuracy of each dataset of PSB under a single feature, and the one obtained with our method are compared in the condition of SB6, as shown in Fig. 13. It can be seen that the classification accuracy of individual datasets under individual features is higher, but not significantly more so than that of our method. On the contrary, most datasets perform much better under our method. This shows that the features selected in this paper are complementary and the feature combination is effective.

In the papers of [17] and [18], seven features are combined to perform segmentation experiments. In addition to the three features used in this paper, they also used four other features includ-

Table 5
Labelling accuracy on princeton benchmark (SB6/SB12).

	SVM SB6 (%)	JointBoost SB6 (%)	ToG [2010] SB6 (%)	ToG [2013] SB6 (%)	ToG [2015] SB6 (%)	Ours SB6 (%)	ToG [2010] SB12 (%)	ToG [2013] SB12 (%)	ToG [2015] SB12 (%)	Ours SB12 (%)
Cup	94.11	93.12	99.1	97.5	99.49	99.59	99.60	99.60	99.73	99.74
Glasses	95.92	93.59	96.10	–	96.78	97.15	97.20	–	97.60	97.79
Airplane	80.43	91.16	95.50	–	95.56	93.90	96.10	–	96.67	95.30
Chair	81.38	95.67	97.80	97.90	97.93	97.05	98.40	99.60	98.67	98.26
Octopus	97.04	96.26	98.60	–	98.61	98.67	98.40	–	98.79	98.80
Table	90.16	97.37	99.10	99.60	99.11	99.25	99.30	99.60	99.55	99.41
Teddy	91.01	85.68	93.30	–	98.00	98.04	98.10	–	98.24	98.27
Plier	92.04	81.55	94.30	–	95.01	95.71	96.20	–	96.22	96.55
Fish	87.05	90.76	95.60	–	96.22	95.63	95.60	–	95.64	95.76
Bird	81.49	81.80	84.20	–	87.51	89.03	87.90	–	88.35	89.48
Mech	81.87	75.73	88.90	90.20	90.56	91.72	90.50	91.30	95.60	96.05
Average	88.41	89.34	94.77	96.30	95.89	95.98	96.12	97.53	96.82	96.85

ToG [2010] is short for Kalogerakis et al. [2010], ToG [2013] is short for Wang et al. [2013], ToG [2015] is short for Guo et al. [2015].




Fig. 14. The comparison between segmentation accuracy under three features and under seven features.

ing curvature (CUR) [55], PCA feature (PCA) [56], shape diameter function (SDF) [57] and distance from medial surface (DIS) [10]. In this paper, several datasets of PSB are randomly selected to perform experiments with the combination of seven features. Under the same experimental conditions, the comparison results with the combination of three features are presented in Fig. 14. The experiment tells us that for most datasets, the combination of three features brings higher classification accuracy than seven features, and for the remaining few datasets, the classification accuracy of the two is very close to each other. This indicates that the three features used in this paper can not only better complement, but also are more suitable for the network structure. In addition, we also trained the combining three features in one SFCN network. However, it didn't produce a higher accuracy result and nearly doubled the training time. So, as a trade-off between performance and efficiency, we empirically trained 3 features separately to perform a segmentation task.

Skip architecture. During the SFCN training process, we utilised skip architecture of five layers and integrated the cross layer information with different steps to improve the SFCN segmentation results, and gradually improve the segmentation details. To verify the skip architecture can do mesh labelling learning and improve the segmentation results, we visualised the cross-layer prediction results with different steps. The segmentation results of several shapes crossing one to five layers are shown in Fig. 15, where (f) is the groundtruth of the corresponding shape. Through tracking of the network computing process, we find that the network convergence is faster with the increase of the cross layers. In addition, it can be seen from the comparison results of cross layers




Fig. 15. The segmentation results of layer information with different steps. (a)–(e) Segmentation results Crossing 1–5 layers; (f) Groundtruth.




Fig. 16. Comparison results before and after optimisation. (a) Segmentation results before optimisation; (b) segmentation results after optimisation. The number below each shape is the segmentation accuracy.

and groundtruth in Fig. 15, that with the increase of cross layers, the classification quality is gradually improved.

Comparison before and after optimisation. In this paper, we use the multi-label graph cut optimisation method [14–16] to optimise the segmentation results of testing shapes, based on the complementarity between features. The comparison results before and after optimisation of several shapes are shown in Fig. 16. As shown in Fig. 16(a), the results before optimisation are the final ones obtained by voting on the three different features tested by SFCN architecture. Fig. 16(b) represents results after optimisation. As the results before optimisation are predicted in the triangle level, the

boundary may be not smooth or there may be prediction errors of individual triangles in some areas. The above problems are well addressed after the optimisation with the multi-label graph cuts optimisation method [14–16], which proves that optimisation plays a significant role. In addition, the number below each figure is the classification accuracy of the corresponding shape, which shows that optimisation can improve classification accuracy.

Limitations. Although the proposed method is effective at completing the 3D shape segmentation task, it has some limitations. Firstly, the shapes involved in the calculation must be manifold meshes, for they can easily determine the connection between triangles. Secondly, the designed SFCN architecture has no feature selection function, thus we carry on the separating training for the three features. Thirdly, to obtain better parameter sharing, the SFCN structure may need all training meshes to be the same triangulation granularity. Therefore, in future work, we will seek to improve the SFCN architecture, making it possible for automatic feature selection, and build an end-to-end structure.

6. Conclusions

We design a novel shape fully convolutional network architecture, which can automatically carry out triangles-to-triangles learning and prediction, and complete the segmentation task with high quality. In the SFCN architecture proposed here, similar to convolution and pooling operation on images, we design a novel shape convolution and pooling operation with a 3D shape represented as a graph structure. Moreover, based on the original image segmentation FCN architecture [10], we first design and implement a new generating operation, which functions as a bridge to facilitate the execution of shape convolution and pooling operations directly on 3D shapes. Additionally, accurate and detailed segmentation of 3D shapes is completed through skip architecture. To produce more accurate segmentation results, we optimise the segmentation results obtained by SFCN prediction by utilising the complementarity between the three geometric features and the multi-label graph cut method [14–16], which can also improve the local smoothness of triangle labels. The experiments show that the proposed method can not only obtain good segmentation results both in large datasets and small ones with the use of a small number of features, but also outperform some existing state-of-the-art shape segmentation methods. More importantly, our method can effectively learn and predict mixed shape datasets either of similar or of different characters, and achieve excellent segmentation results, which demonstrates that our method has strong generalisation ability.

In the future, we want to strengthen our method to overcome some limitations mentioned above and produce more results based on various challenging datasets, such as ShapeNet. Additionally, we hope that the SFCN architecture proposed can be applied to other shape fields, such as shape synthesis, line drawing extraction and so on.

Acknowledgements

We would like to thank all anonymous reviewers for their constructive comments. This research has been supported by the National Science Foundation of China (61321491, 61100110, and 61272219) and the Science and Technology Program of Jiangsu Province (BY2012190 and BY2013072-04).

References

- [1] Yu Y, Zhou K, Xu D, Shi X, Bao H, Guo B, et al. Mesh editing with poisson-based gradient field manipulation. *ACM Trans Graph* 2004;23(3):644–51.
- [2] Yang Y, Xu W, Guo X, Zhou K, Guo B. Boundary-aware multi-domain subspace deformation. *IEEE Trans Vis Comput Graph* 2013;19(19):1633–45.
- [3] Chen X, Zhou B, Lu F, Tan P, Bi L, Tan P. Garment modeling with a depth camera. *ACM Trans Graph* 2015;34(6):1–12.
- [4] Krizhevsky A, Sutskever I, Hinton GE. ImageNet classification with deep convolutional neural networks. *Adv Neural Inf Process Syst* 2012;25(2):2012.
- [5] Szegedy C, Liu W, Jia Y, Sermanet P, Reed S, Anguelov D, et al. Going deeper with convolutions. In: *Proceedings of the 2015 IEEE conference on computer vision and pattern recognition*; 2015. p. 1–9.
- [6] Simonyan K, Zisserman A. Very deep convolutional networks for large-scale image recognition. In: *Proceedings of the 2015 international conference on Learning Representations, ICLR*; 2015.
- [7] Ciresan D, Giusti A, Gambardella LM, Schmidhuber J. Deep neural networks segment neuronal membranes in electron microscopy images. In: *Proceedings of the 2012 advances in neural information processing systems*; 2012. p. 2843–51.
- [8] Farabet C, Couprie C, Najman L, Lecun Y. Learning hierarchical features for scene labeling. *IEEE Trans Pattern Anal Mach Intell* 2013;35(8):1915–29.
- [9] Pinheiro PHO, Collobert R. Recurrent convolutional neural networks for scene parsing. In: *Proceedings of the 2014 international conference on international conference on machine learning*; 2014.
- [10] Long J, Shelhamer E, Darrell T. Fully convolutional networks for semantic segmentation. In: *Proceedings of the 2015 conference on computer vision and pattern recognition, CVPR*; 2015.
- [11] Edwards M, Xie X. Graph based convolutional neural network. In: *Proceedings of the 2016 British machine vision conference, BMVC*; 2016.
- [12] Niepert M, Ahmed M, Kutzkov K. Learning convolutional neural networks for graphs. In: *Proceedings of the 2016 international conference on machine learning, ICML*; 2016.
- [13] Defferrard M, Bresson X, Vandergheynst P. Convolutional neural networks on graphs with fast localized spectral filtering. In: *Proceedings of the 2016 conference on neural information processing systems, NIPS*; 2016.
- [14] Boykov Y, Veksler O, Zabih R. Efficient approximate energy minimization via graph cuts. *IEEE Trans Pattern Anal Mach Intell* 2001;20(12):1222–39.
- [15] Kolmogorov V, Zabih R. What energy functions can be minimized via graph cuts? *IEEE Trans Pattern Anal Mach Intell* 2004;26(2):147–59.
- [16] Boykov Y, Kolmogorov V. An experimental comparison of min-cut/max-flow algorithms for energy minimization in vision. *IEEE Trans Pattern Anal Mach Intell* 2004;26(9):1124–37.
- [17] Kalogerakis E, Hertzmann A, Singh K. Learning 3D mesh segmentation and labeling. *ACM Trans Graph* 2010;29(4):157–66.
- [18] Guo K, Zou D, Chen X. 3D mesh labeling via deep convolutional neural networks. *ACM Trans Graph* 2015;35(1):3.
- [19] Xie Z, Xu K, Liu L, Xiong Y. 3D shape segmentation and labeling via extreme learning machine. *Comput Graph Forum* 2014;33(5):85–95.
- [20] Wang Y, Asafiy S, van Kaick O, Zhang H, Cohen-Or D, Chen B. Active co-analysis of a set of shapes. *ACM Trans Graph* 2012;31(6):157:1–157:10.
- [21] Xie S, Tu Z. Holistically-nested edge detection. In: *Proceedings of the 2015 IEEE international conference on computer vision*; 2015. p. 1395–403.
- [22] Liu F, Shen C, Lin G, Reid I. Learning depth from single monocular images using deep convolutional neural fields. *IEEE Trans Pattern Anal Mach Intell* 2015;38(10):1.
- [23] Dosovitskiy A, Fischer P, Ilg E, Häusser P, Hazrbaş C, Golkov V, et al. FlowNet: learning optical flow with convolutional networks. In: *Proceedings of the 2015 IEEE international conference on computer vision*; 2015. p. 2758–66.
- [24] Simo-Serra E, Lizuka S, Sasaki K, Ishikawa H. Learning to simplify: fully convolutional networks for rough sketch cleanup. *ACM Trans Graph* 2016;35(4):1–11.
- [25] Bruna J, Zaremba W, Szlam A, LeCun Y. Spectral networks and locally connected networks on graphs. In: *Proceedings of the 2014 international conference on languages and arts, ICLA*; 2014.
- [26] Duvenaudy D, Maclaurin D, Aguilera-Iparraguirre J, Gomez-Bombarelli R, Hirzel T, Aspuru-Guzik A, et al. Convolutional networks on graphs for learning molecular fingerprints. In: *Proceedings of the advances in neural information processing systems*; 2015. p. 2224–32.
- [27] Wu Z, Song S, Khosla A, Yu F, Zhang L, Tang X, et al. A deep representation for volumetric shapes. In: *Proceedings of the 2015 IEEE conference on computer vision and pattern recognition*; 2015.
- [28] Qi CR, Su H, Niener M, Dai A, Yan M, Guibas L. Volumetric and multi-view CNNs for object classification on 3d data. In: *Proceedings of the 2016 computer vision and pattern recognition*; 2016.
- [29] Kalogerakis E, Averkiou M, Maji S, Chaudhuri S. 3D shape segmentation with projective convolutional networks. In: *Proceedings of the 2017 IEEE computer vision and pattern recognition*; 2017.
- [30] Yi L, Su H, Guo X, Guibas L. SyncSpecCNN: synchronized spectral CNN for 3D shape segmentation. In: *Proceedings of the 2017 IEEE computer vision and pattern recognition*; 2017.
- [31] Su H, Qi C, Mo K, Guibas L. PointNet: deep learning on point sets for 3D classification and segmentation. In: *Proceedings of the 2017 IEEE computer vision and pattern recognition*; 2017.
- [32] Masci J, Boscaini D, Bronstein MM, Vandergheynst P. Geodesic convolutional neural networks on Riemannian manifolds. In: *Proceedings of the 2015 international IEEE workshop on 3D representation for recognition, 3DRR*; 2015.
- [33] Boscaini D, Masci J, Rodol E, Bronstein MM. Learning shape correspondence with anisotropic convolutional neural networks. In: *Proceedings of the 2016 conference on neural information processing systems, NIPS*; 2016.
- [34] Monti F, Boscaini D, Masci J, Rodol E, Svoboda J, Bronstein MM. Geometric deep learning on graphs and manifolds using mixture model CNNs. In: *Pro-*

- ceedings of the 2017 conference on computer vision and pattern recognition, CVPR; 2017.
- [35] Maron H, Galun M, Aigerman N, Trope M, Dym N, Yumer E, et al. Convolutional neural networks on surfaces via seamless toric covers. In: Proceedings of the 2017 SIGGRAPH; 2017.
 - [36] Boscaini D, Masci J, Bronstein MM, Castellani U. Learning class-specific descriptors for deformable shapes using localized spectral convolutional networks. *Eurographics Symp Geom Process* 2015;34(5):13–23.
 - [37] Litany O, Remez T, Rodol E, Bronstein AM, Bronstein MM. Deep functional maps: structured prediction for dense shape correspondence. In: Proceedings of the 2017 international conference on computer vision, ICCV; 2017.
 - [38] Bronstein MM, Bruna J, LeCun Y, Szlam A, Vandergheynst P. Geometric Deep Learning: Going beyond Euclidean data. *IEEE Signal Processing Magazine* 2016;34(4):18–42.
 - [39] Wang Y, Gongy M, Wang T, Cohen-Or D, Zhang H, Chen B. Projective analysis for 3D shape segmentation. *ACM Trans Graph* 2013;32(6):1–12.
 - [40] Huang Q, Koltun V, Guibas L. Joint shape segmentation with linear programming. *ACM Trans Graph (TOG)* 2011;30(6) 125:1–125:12.
 - [41] Sidi O, van Kaick O, Kleiman Y, Zhang H, Cohen-Or D. Unsupervised co-segmentation of a set of shapes via descriptor-space spectral clustering. In: Proceedings of the 2011 SIGGRAPH Asia conference; 2011. p. 1.
 - [42] Hu R, Fan L, Liu L. Co-segmentation of 3D shapes via subspace clustering. *Comput Graph Forum* 2012;31(5):1703–13.
 - [43] Meng M, Xia J, Luo J, He Y. Unsupervised co-segmentation for 3D shapes using iterative multi-label optimization. *Comput-Aided Des* 2013;45(2):312–20.
 - [44] Xu K, Li H, Zhang H, Cohen-Or D, Xiong Y, Cheng Z-Q. Style-content separation by anisotropic part scales. *ACM Trans Graph* 2010;29(6):184.
 - [45] Kreyov V, Julius D, Sheffer A. Model composition from interchangeable components. In: Proceedings of the conference on computer graphics & applications; 2007. p. 129–38.
 - [46] Dhillon IS, Guan Y, Kulis B. Weighted graph cuts without eigenvectors: a multilevel approach. *IEEE Trans Pattern Anal Mach Intell* 2007;29(11):1944–57.
 - [47] Hilaga M, Shinagawa Y, Kohmura T, Kunii TL. Topology matching for fully automatic similarity estimation of 3d shapes. In: Proceedings of the conference on computer graphics and interactive techniques; 2001. p. 203–12.
 - [48] Belongie S, Malik J, Puzicha J. Shape matching and object recognition using shape contexts. *IEEE Trans Pattern Anal Mach Intell* 2002;24(4):509–22.
 - [49] Johnson AE, Hebert M. Using spin images for efficient object recognition in cluttered 3d scenes. *IEEE Trans Pattern Anal Mach Intell* 1999;21(5):433–49.
 - [50] Jia Y, Shelhamer E, Donahue J, Karayev S, Long J, Girshick R, et al.; 2014. arXiv preprint arXiv:1408.5093.
 - [51] Chen X, Ieksey Golovinskiy, Funkhouser T. A benchmark for 3D mesh segmentation. *ACM Trans Graph* 2009;28(3):341–52.
 - [52] Kim VG, Li W, Mitra NJ, Chaudhuri S, DiVerdi S, Funkhouser T. Learning part-based templates from large collections of 3D shapes. *ACM Trans Graph* 2013;32(4):1.
 - [53] Chang C-C, Lin C-J. LIBSVM: a library for support vector machines. *ACM Trans Intell Syst Technol* 2007;2(3, article 27):389–96.
 - [54] Torralba A, Murphy KP, Freeman WT. Sharing visual features for multi-class and multiview object detection. *IEEE Trans Pattern Anal Mach Intell* 2007;29(5):854–69.
 - [55] Kavukcuoglu K, Ranzato M, Fergus R, LeCun Y. Learning invariant features through topographic filter maps. In: Proceedings of the IEEE conference on computer vision & pattern recognition; 2009. p. 1605–12.
 - [56] Shapira L, Shalom S, Shamir A, Cohen-Or D, Zhang H. Contextual part analogies in 3D objects. *Int J Comput Vis* 2010;89(2):309–26.
 - [57] Liu R, Zhang H, Shamir A, Cohen-Or D. A part-aware surface metric for shape analysis. *Comput Graph Forum* 2009;28(2):397–406.